Python and AI
for Modern Threat Intelligence

Cyber Threat Intelligence

Threat Intelligence Lifecycle

Direction involves identifying critical assets and processes requiring protection while determining specific intelligence needs for the organization.

Feedback gathers input from data consumers to continuously improve the value, accuracy, and actionability of the threat intelligence.

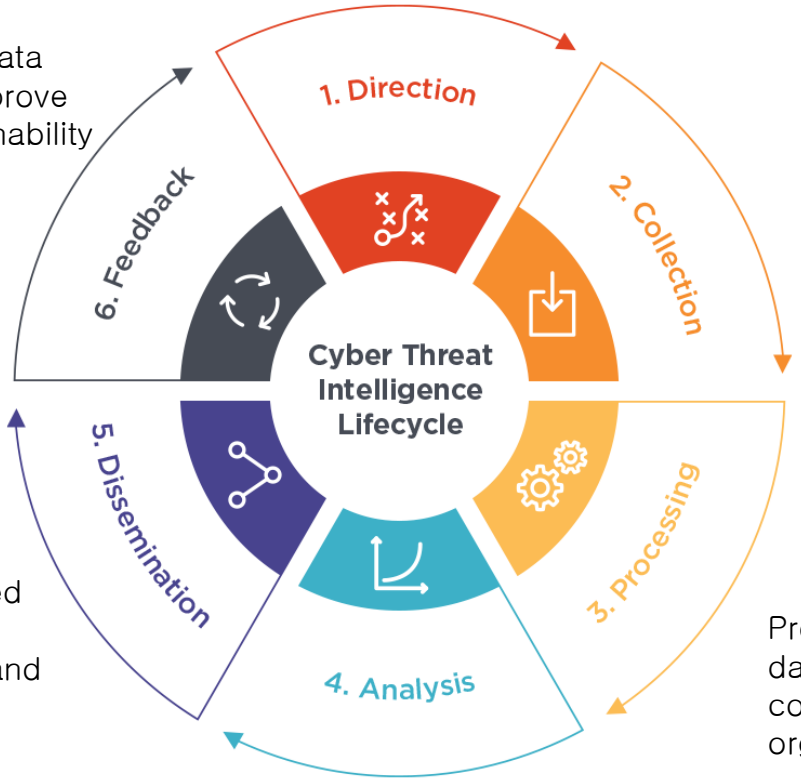

Collection accumulates relevant information from internal and externa sources to address key intelligence requirements.

Dissemination distributes analyzed intelligence to appropriate stakeholders in optimal formats and frequencies.

Processing transforms raw collected data into a format that can be easily consumed throughout the organization.

Analysis converts processed information into actionable intelligence for decision-making with clear findings and recommendations.

https://threat.media/definition/what-is-the-threat-intelligence-lifecycle/

Threat Intelligence with Security Teams

- Security Operation Center
  - Enhances detection and response with threat context
- Incident Response
  - Provides attack patterns and indicators for effective response
- Vulnerability Management
  - Prioritizes patching based on active exploit information
- Malware Analysis
  - Offers context on malware families and associated threat actors
- Business Operations
  - Informs risk-based decisions and resource allocation
- System Engineering & IT
  - Guides security control implementation against specific threats

Threat Intelligence Use Case (SOC)

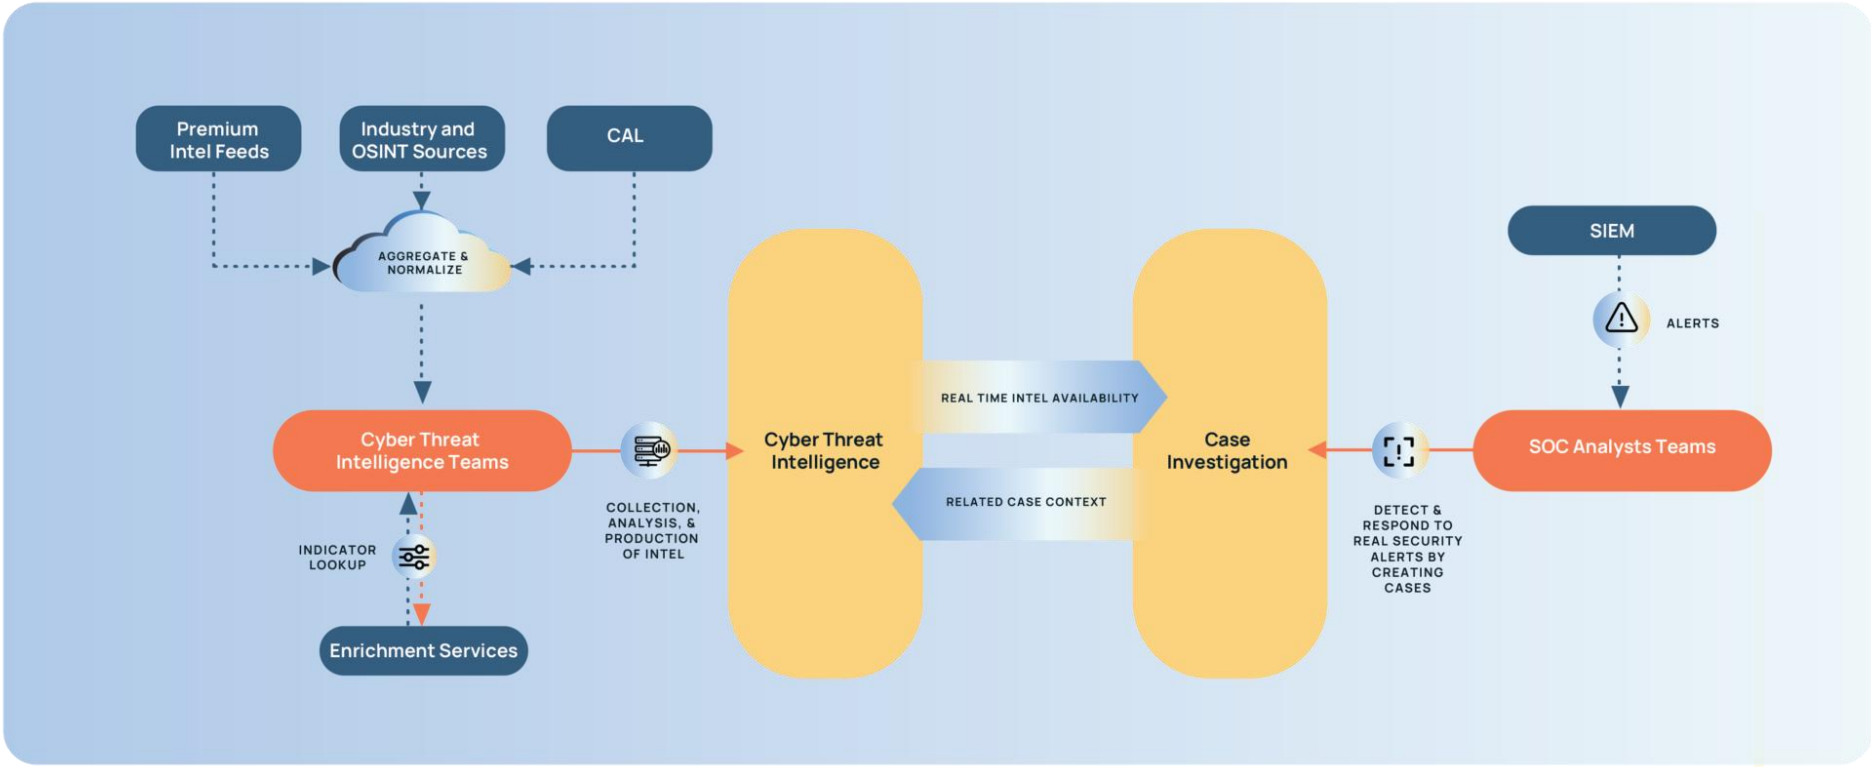

https://threatconnect.com/blog/improve-soc-efficiency-with-intelligence-powered-security-operations/

Type of Threat Intelligence

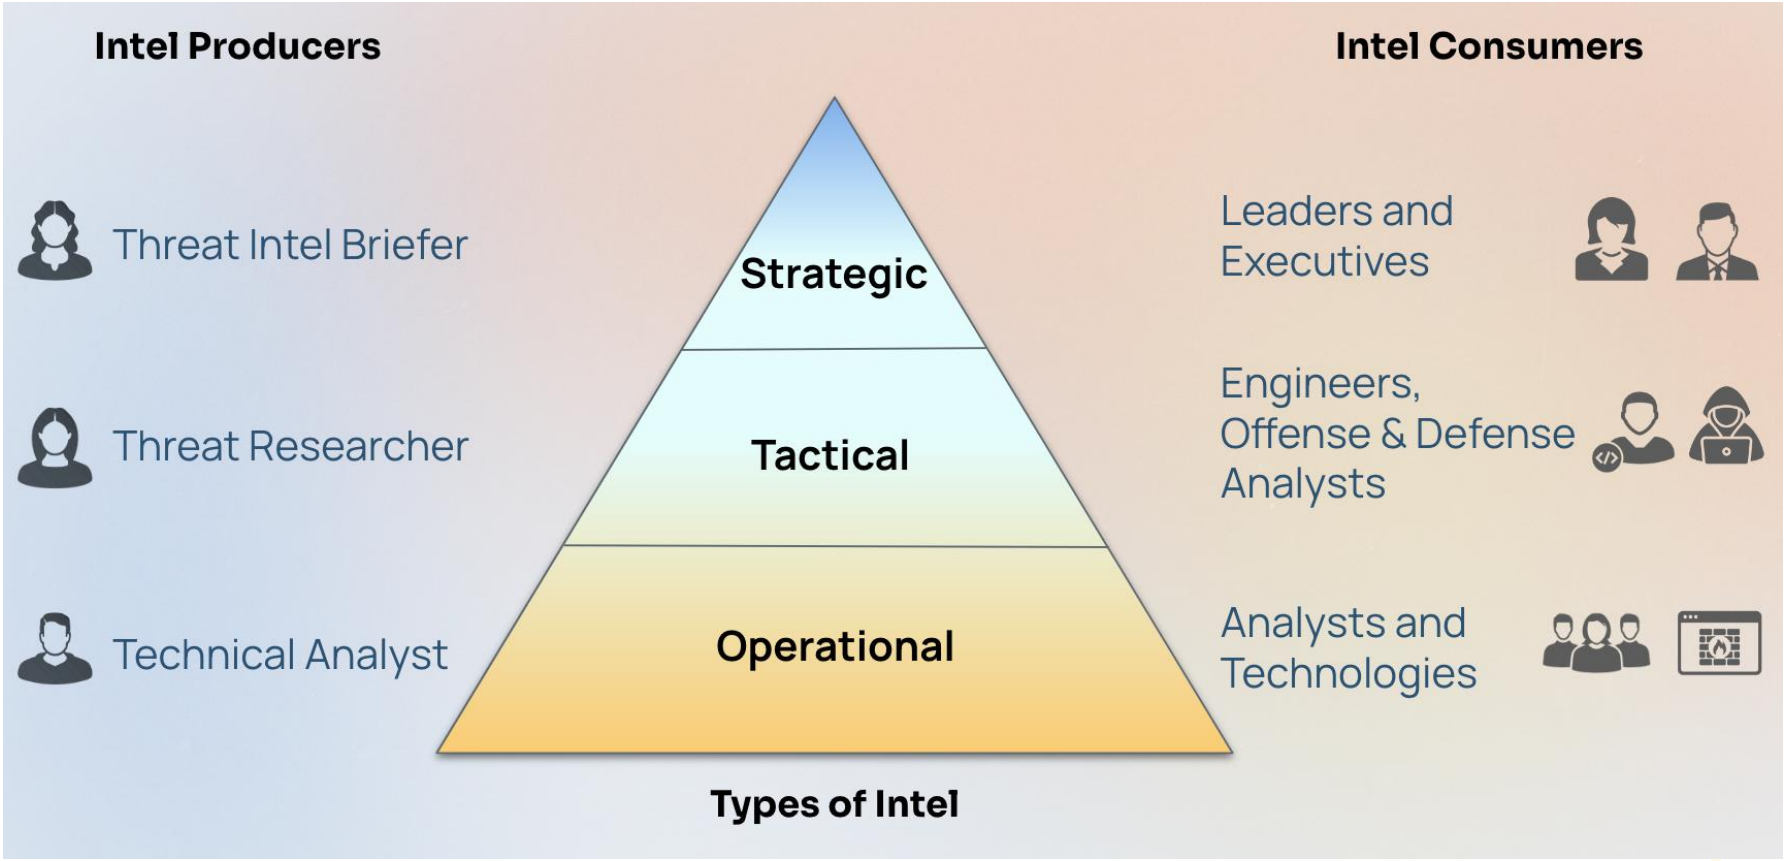

https://threatconnect.com/blog/the-7-tenets-of-threat-intelligence-operations-tenet-6-integrates-and-automates-threat-intel/

The Pyramid of Pain

TTPs describe how threat actors conduct their operations, offering a broader context than IOCs.

IOCs are pieces of data that identify potentially malicious activity on a system or network.

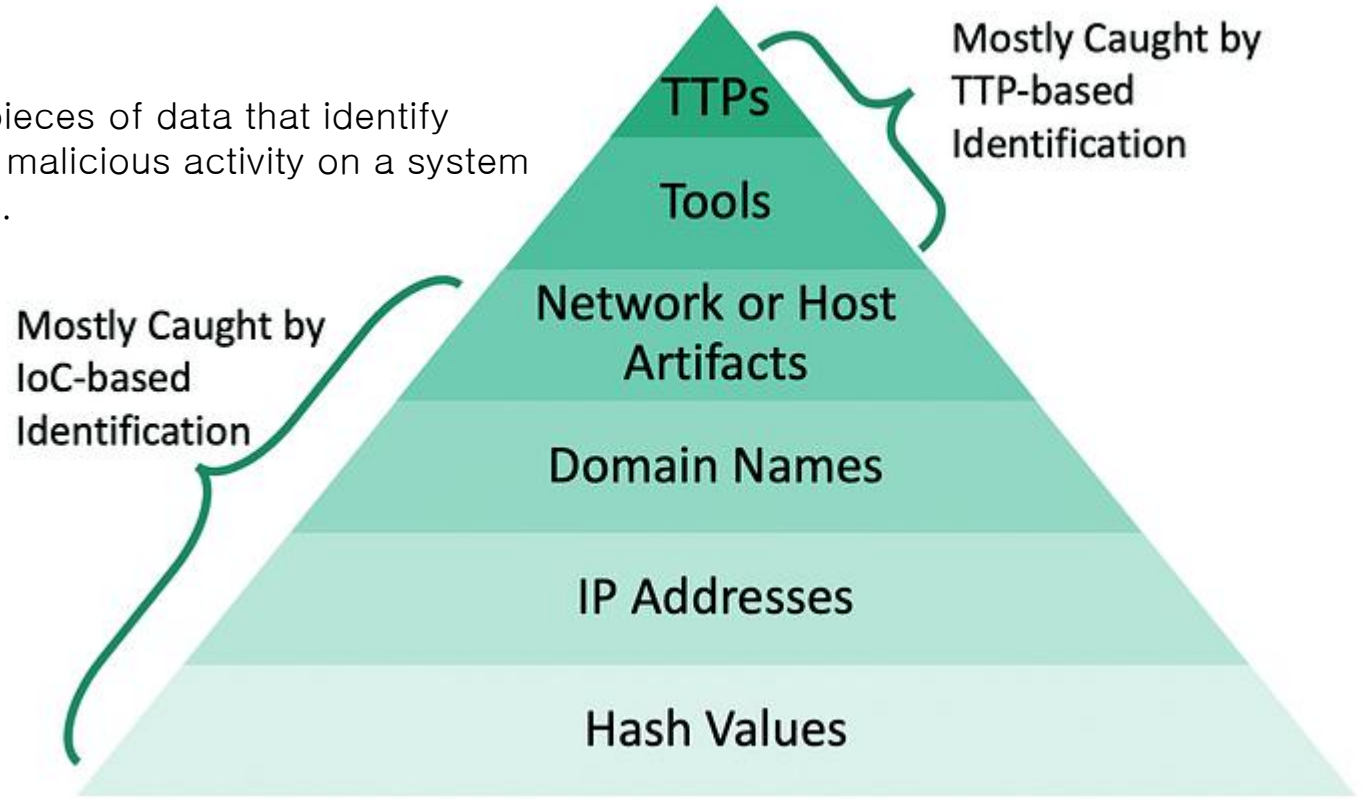

https://ibrahimakkdag.medium.com/tips-for-threat-hunters-comparison-of-indicators-of-compromise-iocs-and-tactics-techniques-and-47bc268fe7ff

All Source Intelligence

Communications intelligence,
Electronic intelligence, …

Interrogations, Agent networks, …

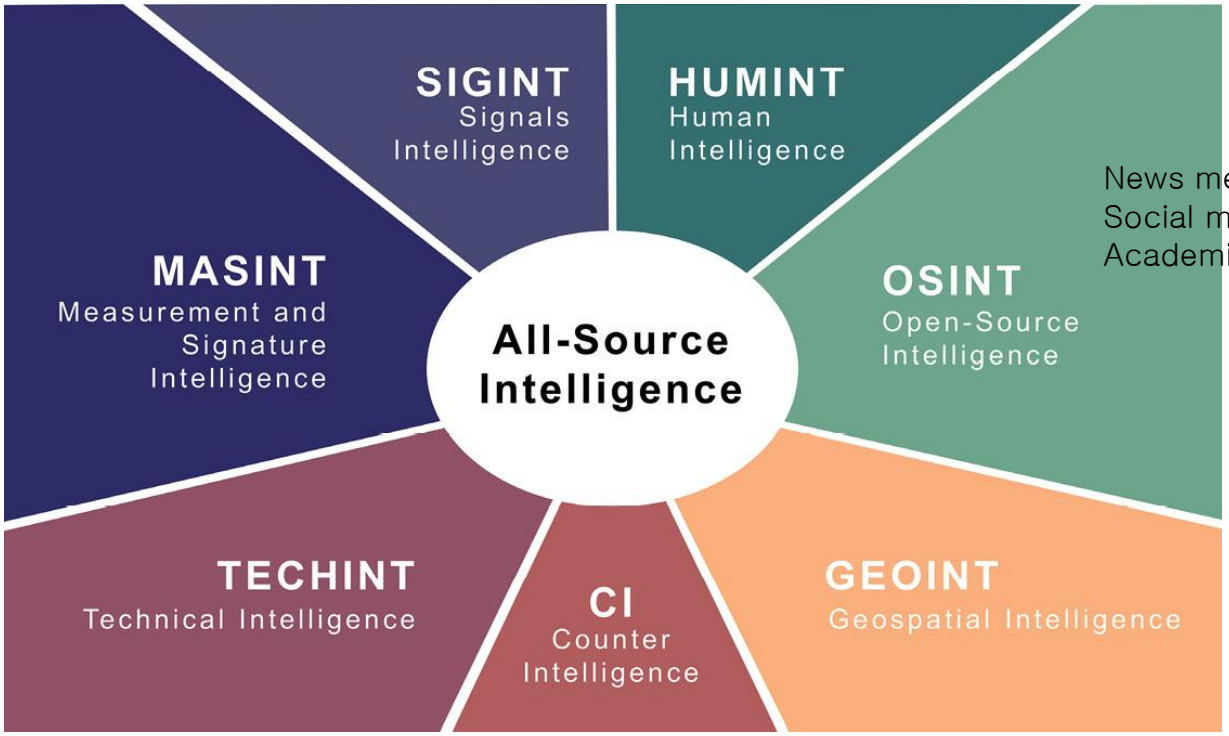

Radar signatures,
Nuclear intelligence,
Materials analysis, …

News media,
Social media,
Academic publications, …

Weapons systems analysis,
Military equipment capabilities,
Technical specifications, …

Double agent operations,
Surveillance detection,
Penetration testing, …

Satellite imagery,
Aerial photography,
Mapping/terrain data, …

https://tjaglcs.army.mil/Periodicals/The-Army-Lawyer/tal-2022-issue-4/Post/4255/No-1-Distinguishing-Between-Operational-and-Intelligence-Activities

Pivoting

- Pivoting involves leveraging initial security discoveries to broaden investigative scope

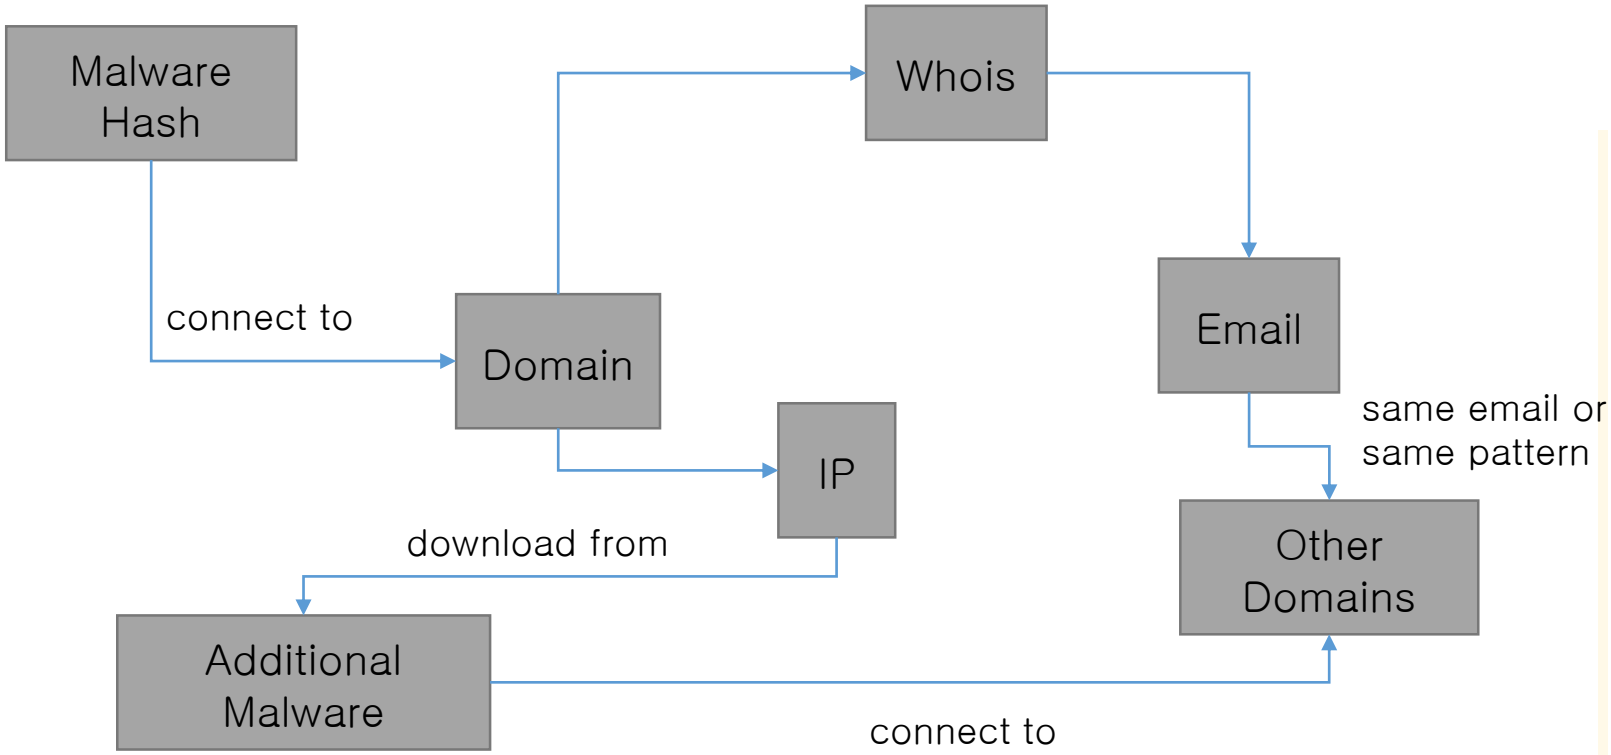

Pivoting

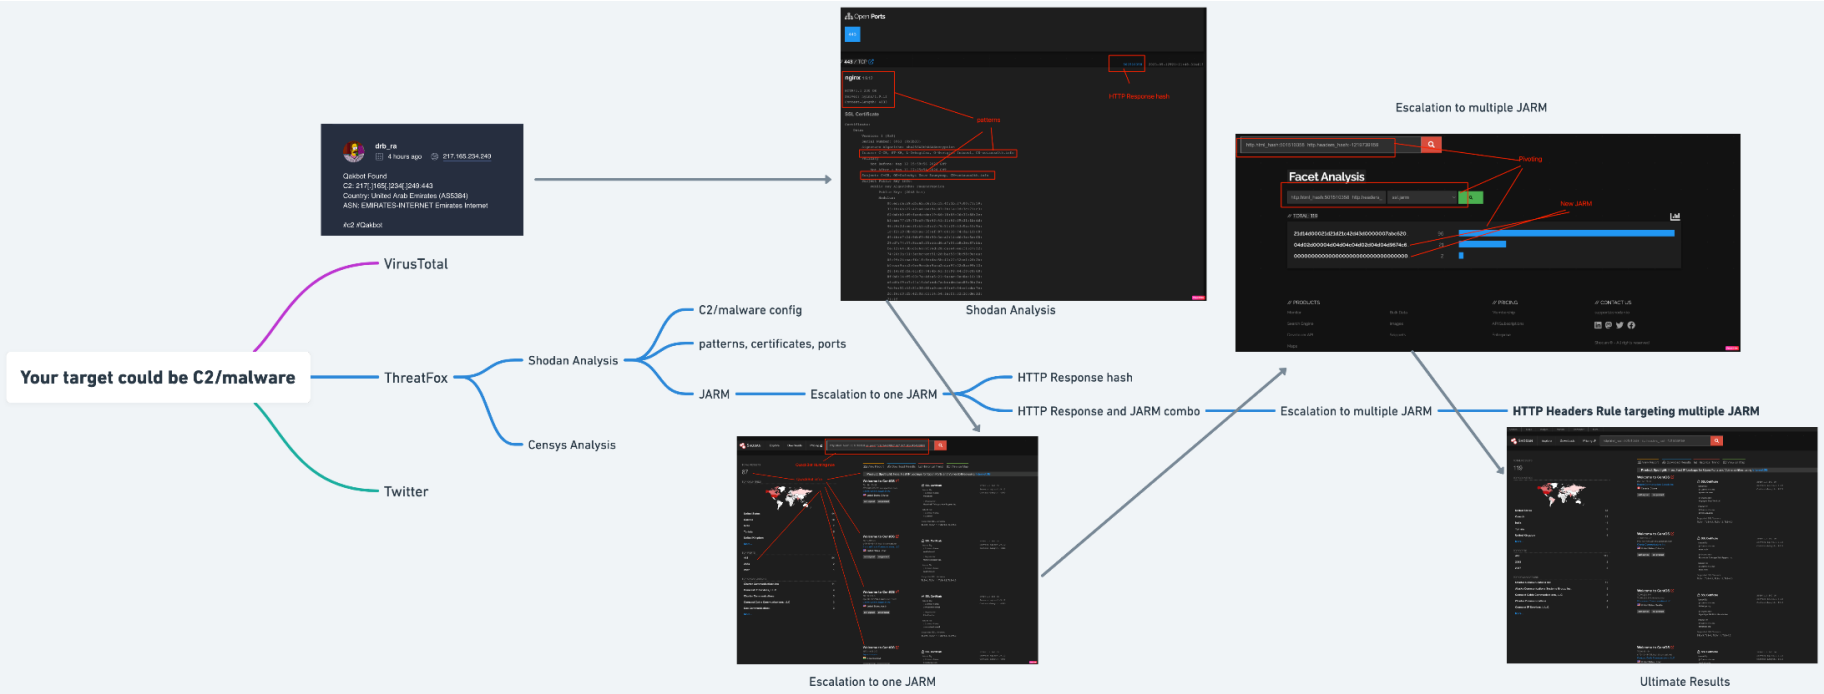

https://detect.fyi/hunting-malicious-infrastructure-using-jarm-and-http-response-bb4a039d4119

Python for Threat Intelligence

Python automates security needs assessment through scripted tools that identify critical assets

Python builds interactive systems for collecting stakeholder feedback on intelligence products

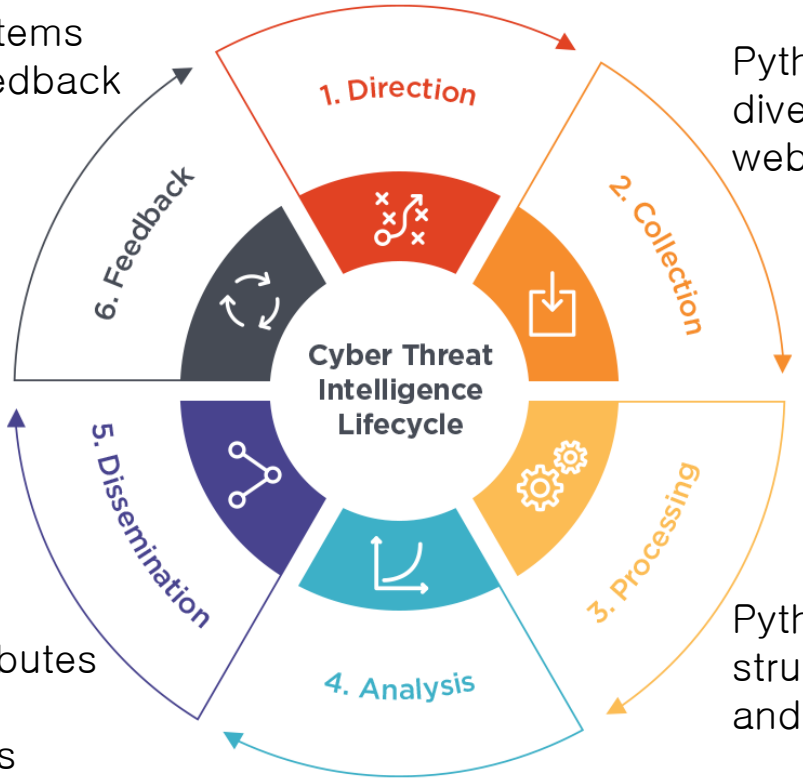

Python collects intelligence from diverse sources using APIs and web scraping libraries

Python transforms raw data into structured formats using Pandas and regex for analysis

Python generates and distributes visual intelligence reports through automated channels

Python analyzes threats with data science libraries to uncover patterns and actionable insights

STIX & TAXII

- STIX (Structured Threat Information Expression)
  - Standardized language for cyber threat intelligence representation
  - Enables organizations to share structured threat information using objects, relationships, and properties
- TAXII (Trusted Automated Exchange of Intelligence Information)
  - Transport mechanism for sharing cyber threat intelligence
  - Defines APIs for STIX content exchange between producers and consumers of threat information

STIX Scenario

- STIX Producer: Company A
  - Creates Indicator SDO with CryptoLocker hash pattern and Malware SDO with malware details
  - Links objects with "indicates" Relationship SRO and shares bundle via TAXII server
- STIX Consumer: Company B
  - Receives Company A's intelligence and finds matching malware on their network
  - Creates Sighting SRO to report detection and publishes back to community via TAXII

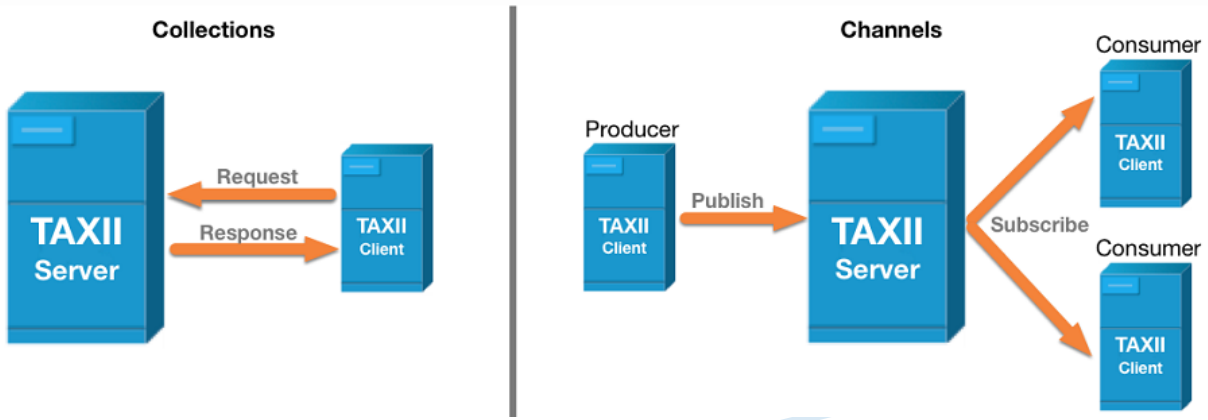

STIX Scenario

- Contains detection info: name, pattern, type, validity timeframe
- Company A example: CryptoLocker SHA-256 hash marked as "malicious-activity"

Indicator Object


```
{
   "type": "indicator",
   "spec_version": "2.1",
   "id": "indicator--71312c48-925d-44b7-b10e-c11086995358",
   "created": "2017-02-06T09:13:07.243000Z",
   "modified": "2017-02-06T09:13:07.243000Z",
   "name": "CryptoLocker Hash",
   "description": "This file is a part of CryptoLocker",
   "pattern": "[file:hashes.'SHA-256' = '46afeb295883a5efd6639d4197eb18bcba3bff49125b810ca4b9509b9ce4dfbf']",
   "pattern_type": "stix",
   "indicator_types": ["malicious-activity"],
   "valid_from": "2017-01-01T09:00:00.000000Z"
}
```

STIX Scenario

- Contains required properties: type, spec_version, id, created
- Provides descriptive context about the malware

Malware Object

```
{
    "type": "malware",
    "id": "malware--81be4588-96a8-4de2-9938-9e16130ce7e6",
    "spec_version": "2.1",
    "created": "2017-02-06T09:26:21.647000Z",
    "modified": "2017-02-06T09:26:21.647000Z",
    "name": "CryptoLocker",
    "description": "CryptoLocker is known to hold files hostage for ransom.",
    "malware_types": ["ransomware"]
}
```

STIX Scenario

- Links Indicator to Malware SDOs using source_ref, target_ref, and relationship_type
- Company A example: Indicator "indicates" Malware, connecting detection pattern to CryptoLocker

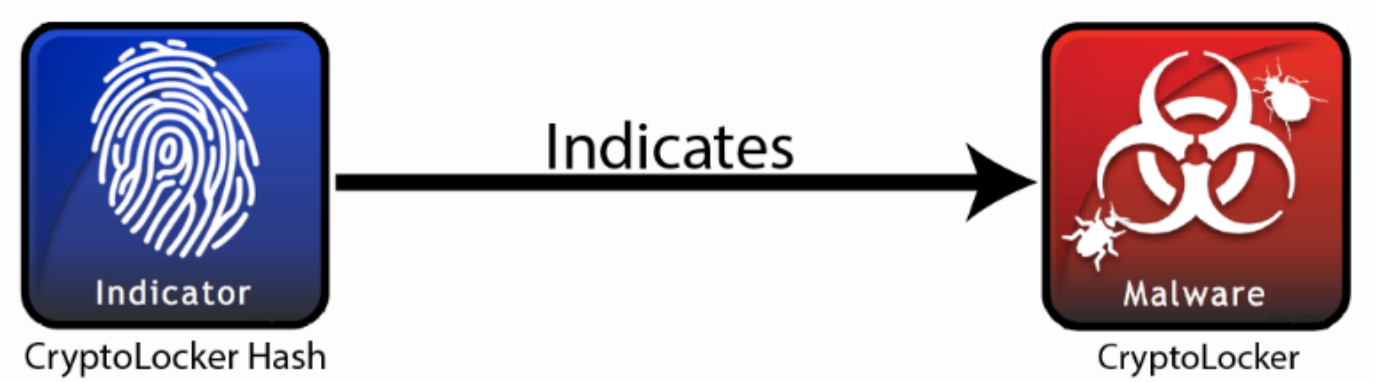

STIX Scenario

- Reports when an organization observes a previously shared threat indicator
- Company B example: Documents detection of Company A's CryptoLocker hash on their network

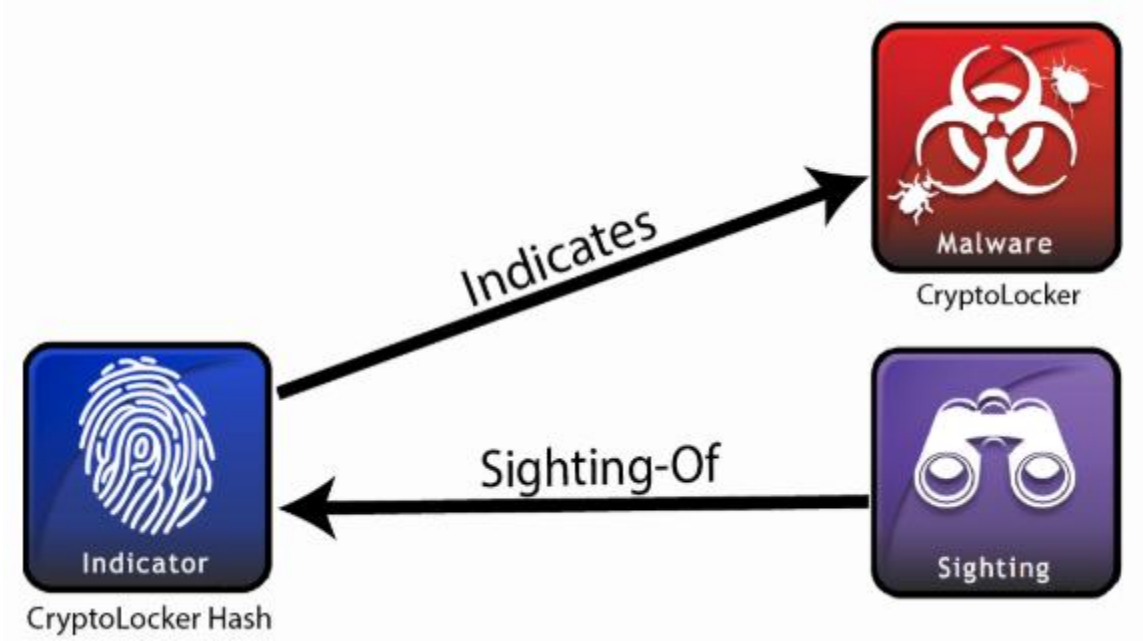

STIX Scenario

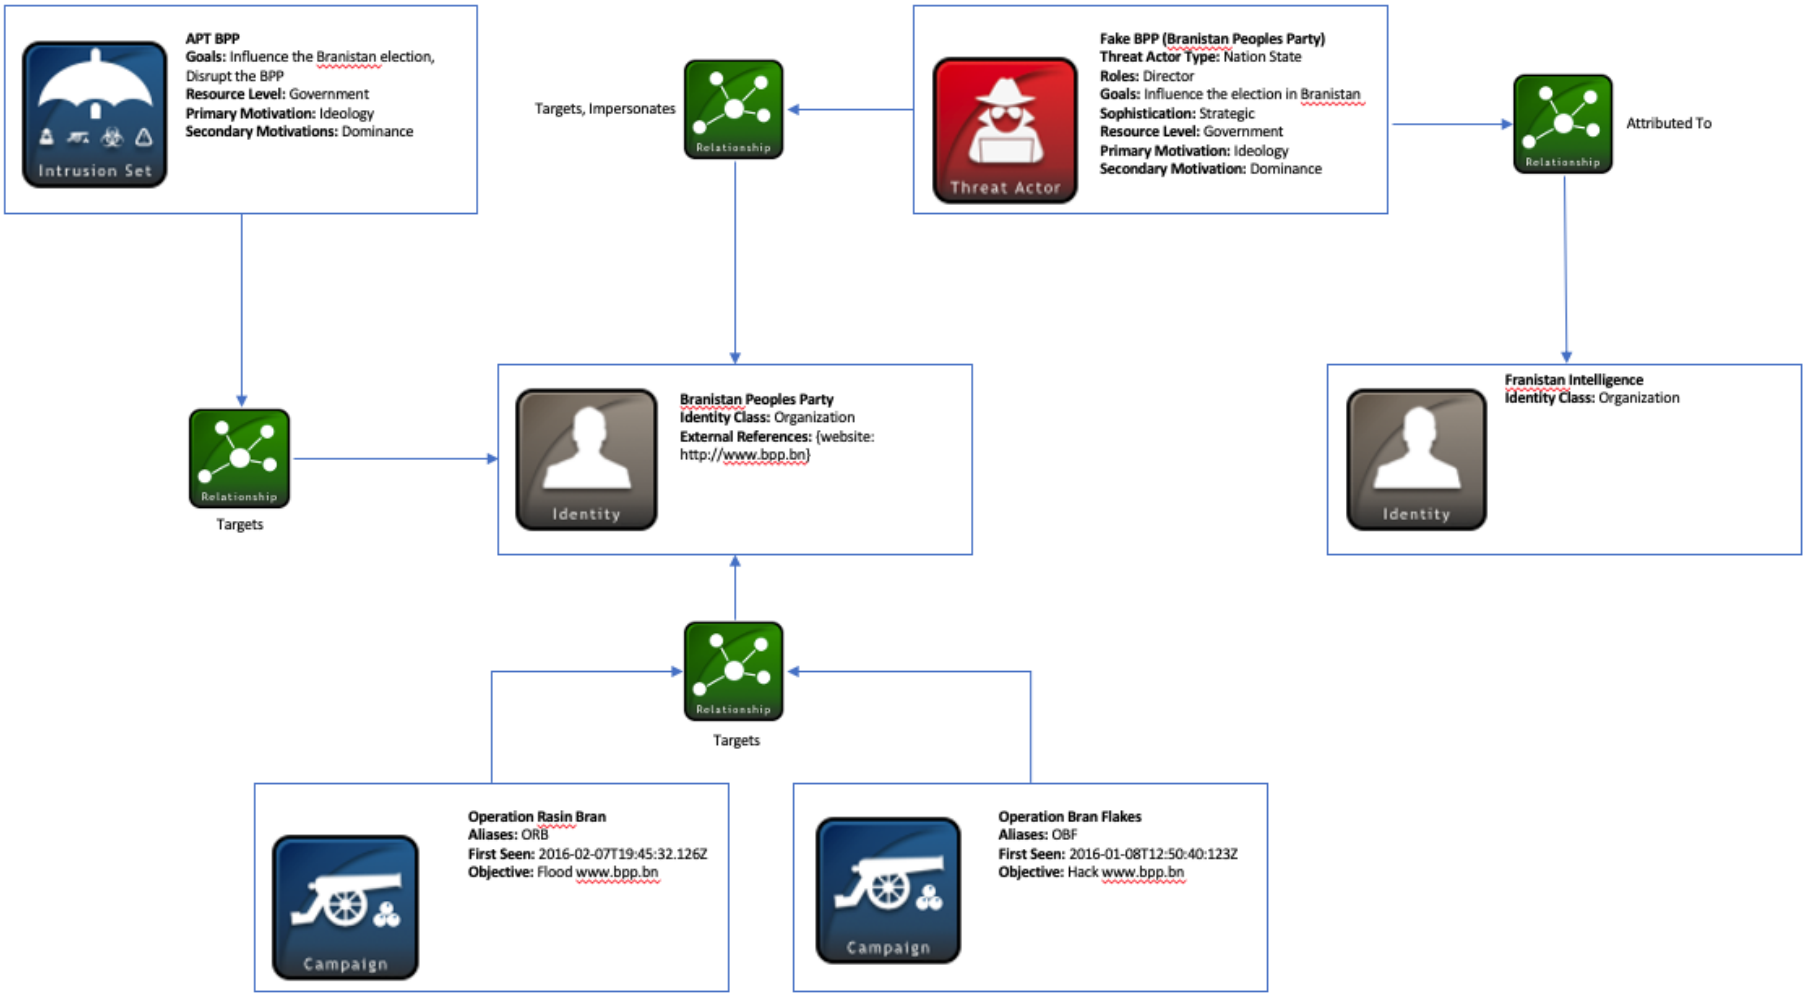

https://oasis-open.github.io/cti-documentation/examples/defining-campaign-ta-is

MITRE ATT&CK

- Free knowledge base of real-world adversary tactics and techniques
- Foundation for threat modeling across government and industry
- ATT&CK data is in STIX format for easy sharing and analysis

MITRE ATT&CK

ATT&CK concept	STIX object type	Custom type?
Matrix	`x-mitre-matrix`	yes
Tactic	`x-mitre-tactic`	yes
Technique	attack-pattern	no
Sub-technique	attack-pattern where `x_mitre_is_subtechnique = true`	no
Procedure	relationship where `relationship_type = "uses"` and `target_ref` is an `attack-pattern`	no
Mitigation	course-of-action	no
Group	intrusion-set	no
Software	malware or tool	no
Collection[1]	`x-mitre-collection`	yes
Data Source	`x-mitre-data-source`	yes
Campaign	campaign	no
Asset	`x-mitre-asset`	yes

https://github.com/mitre-attack/attack-stix-data/blob/master/USAGE.md

MITRE ATT&CK

Source Type	Relationship Type	Target Type	Custom Type?	About
`intrusion-set`	`uses`	`malware` or `tool`	No	Group using a software.
`intrusion-set`	`uses`	`attack-pattern`	No	Group using a technique, which is also considered a procedure example.
`malware` or `tool`	`uses`	`attack-pattern`	No	Software using a technique, which is also considered a procedure example.
`campaign`	`uses`	`malware` or `tool`	No	Campaign using a software.
`campaign`	`uses`	`attack-pattern`	No	Campaign using a technique, which is also considered a procedure example.
`campaign`	`attributed-to`	`intrusion-set`	No	Campaign attributed to a group.
`course-of-action`	`mitigates`	`attack-pattern`	No	Mitigation mitigating technique.
`attack-pattern`	`subtechnique-of`	`attack-pattern`	Yes	Sub-technique of a technique, where the `source_ref` is the sub-technique and the `target_ref` is the parent technique.
`x-mitre-data-component`	`detects`	`attack-pattern`	Yes	Data component detecting a technique.
`attack-pattern`	`targets`	`x-mitre-asset`	Yes	Technique targets an asset.
any type	`revoked-by`	any type	Yes	The target object is a replacement for the source object. Only occurs where the objects are of the same type, and the source object will have the property `revoked = true`. See Working with deprecated and revoked objects for more information on revoked objects.

https://github.com/mitre-attack/attack-stix-data/blob/master/USAGE.md

Data Collection

- API Querying
  - Uses Python scripts to retrieve structured data directly from providers
  - Gathers threat intelligence efficiently via platform APIs
- Web Scraping
  - Extracts unstructured data from websites using tools like BeautifulSoup
  - Collects IOCs from hacker forums and dark web sources
- Open Source Intelligence (OSINT)
  - Automates intelligence collection from feeds and unstructured sources
  - Monitors dark web activities and analyzes threat reports
- Offensive Threat Intelligence
  - Scans and exploits vulnerabilities in attackers' infrastructure
  - Analyzes malware commands to disable threats or infiltrate systems

API Providers

- Free/Community
  - MISP: Open-source threat intelligence platform for sharing, storing, and correlating indicators
  - PhishTank: Community-driven database of verified phishing websites
  - Validin: Historical DNS record lookup service for tracking domain changes
  - Malware Bazaar: Repository for sharing and analyzing malware samples
  - Pulsedive: Threat intelligence platform providing enriched IOC data
  - DNSQuery: DNS lookup tool for analyzing domain information
  - UrlScan: Website scanner that analyzes and detects suspicious websites
  - OTX: AlienVault's Open Threat Exchange for community threat intelligence sharing
- Enterprise
  - VirusTotal: Multi-engine malware scanning and file reputation service
  - Censys: Internet-wide scanning platform for attack surface management
  - Flare: Threat intelligence platform focused on dark web monitoring
  - Shodan: Search engine for internet-connected devices and vulnerabilities
  - MS Defender for TI: Microsoft's threat intelligence offering within Defender suite
  - Recorded Future: AI-powered threat intelligence platform with real-time risk assessment
  - Mandiant Threat Intelligence: Advanced threat intelligence with actor profiling and vulnerability research

Data Processing

- Structure Your Dataset
  - Integrate multiple source formats (CSV, JSON) into a unified, organized collection
  - Ensure field consistency through standardization processes for seamless analysis
- Language Conversion
  - Transform text between languages to enable global threat intelligence operations
  - Maintain consistent encoding protocols across multilingual data to preserve integrity
- Image Text Extraction
  - Convert visual information to text using OCR for documents and screenshots containing threat indicators
- Noise Elimination
  - Remove irrelevant data points that don't contribute meaningful intelligence value
  - Develop filtering parameters based on statistical methods and known benign patterns
- Pandas, NumPy, Elasticsearch, Apache Spark, SQLite, KQL, …

Data Analysis

- Statistics
  - Identify threat behavior patterns through correlation and regression analysis
  - Apply statistical testing to validate security event relationships
- Visualization
  - Create charts and diagrams to represent complex threat data clearly
  - Display geographical attack distributions and malware propagation timelines
- Enrich and Pivot
  - Add contextual information to raw security data from external intelligence sources
  - Examine multiple data attributes to discover hidden attack relationships
- Intelligence Identification
  - Distinguish meaningful threats from routine security alerts
  - Filter out false positives to focus analysis on genuine security risks
- Matplotlib, Bokeh, Pyvis, NetworkX, Seaborn, ggplot, Pygal, …

Disseminating Threat Intelligence

- Definition
  - The delivery of critical security insights to relevant decision-makers and teams.
- Tools
  - VThunting.py: Python tool for automated VirusTotal hunting and intelligence gathering
  - PyMISP: Python library for interacting with MISP threat intelligence platforms
  - PySTIX: Python implementation for working with STIX threat intelligence format
  - Discord: Messaging platform that can be leveraged for real-time threat intel distribution
  - Slack: Collaboration tool with channels and integrations for sharing security alerts
  - Jinja2: Template engine used to generate standardized threat intelligence reports
  - MS Teams: Microsoft's collaboration platform with channels for threat intelligence distribution

Generative AI and Threat Intelligence

LLM

- Definition
  - Large Language Models are advanced AI systems trained on vast amounts of text data to understand, generate, and manipulate human language across diverse tasks and contexts

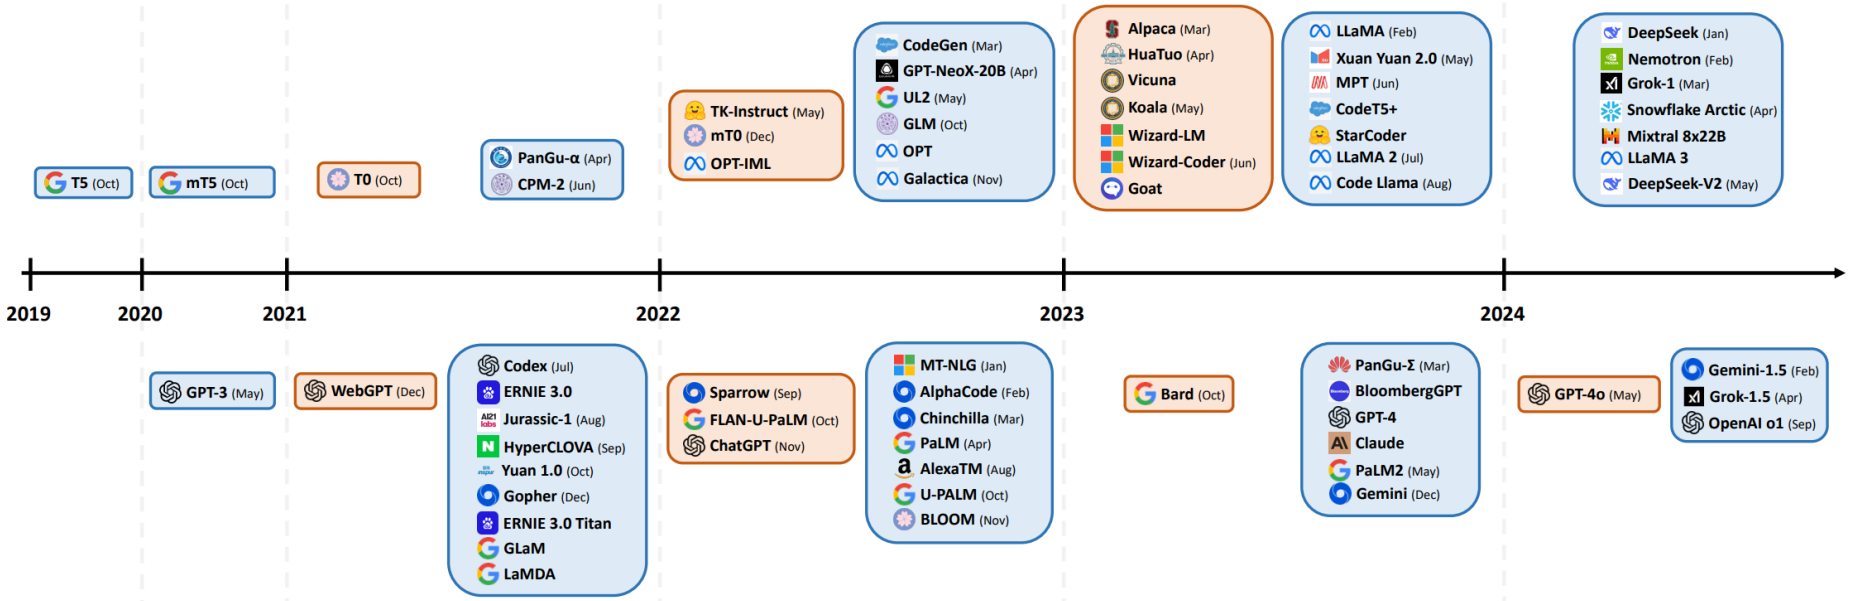

- Blue = pre-trained, Orange = instruction-tuned
- Top = open-source, Bottom = closed-source

Naveed, Humza, et al. "A comprehensive overview of large language models." arXiv preprint arXiv:2307.06435 (2023).

LLM

- Text in, Text out!
  - LLMs were focused solely on processing and generating plain text, mainly used for tasks like summarization, translation, and classification

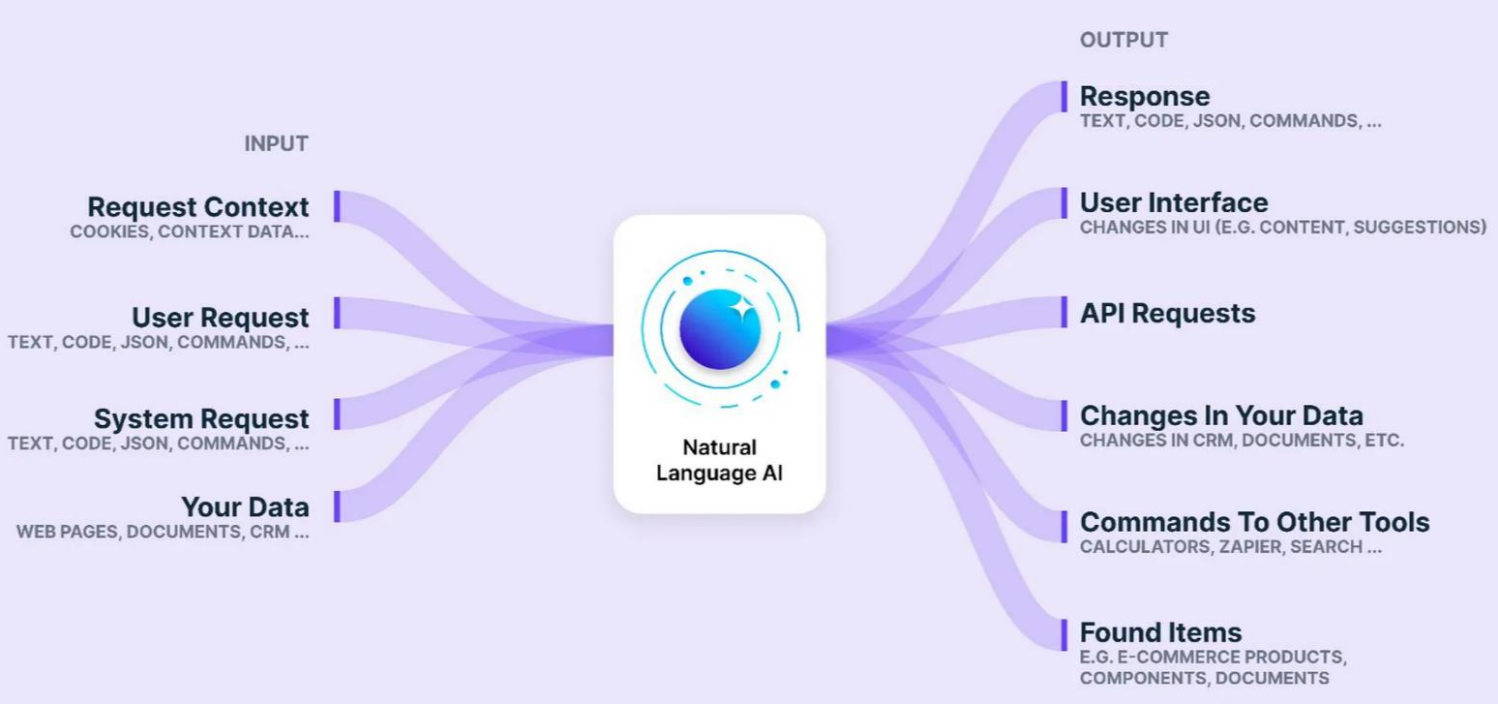

But with the advent of function calling, LLM now have the ability to interact with external tools and APIs—unlocking far more dynamic, real-world applications

LLM

- Anything in, Anything out!

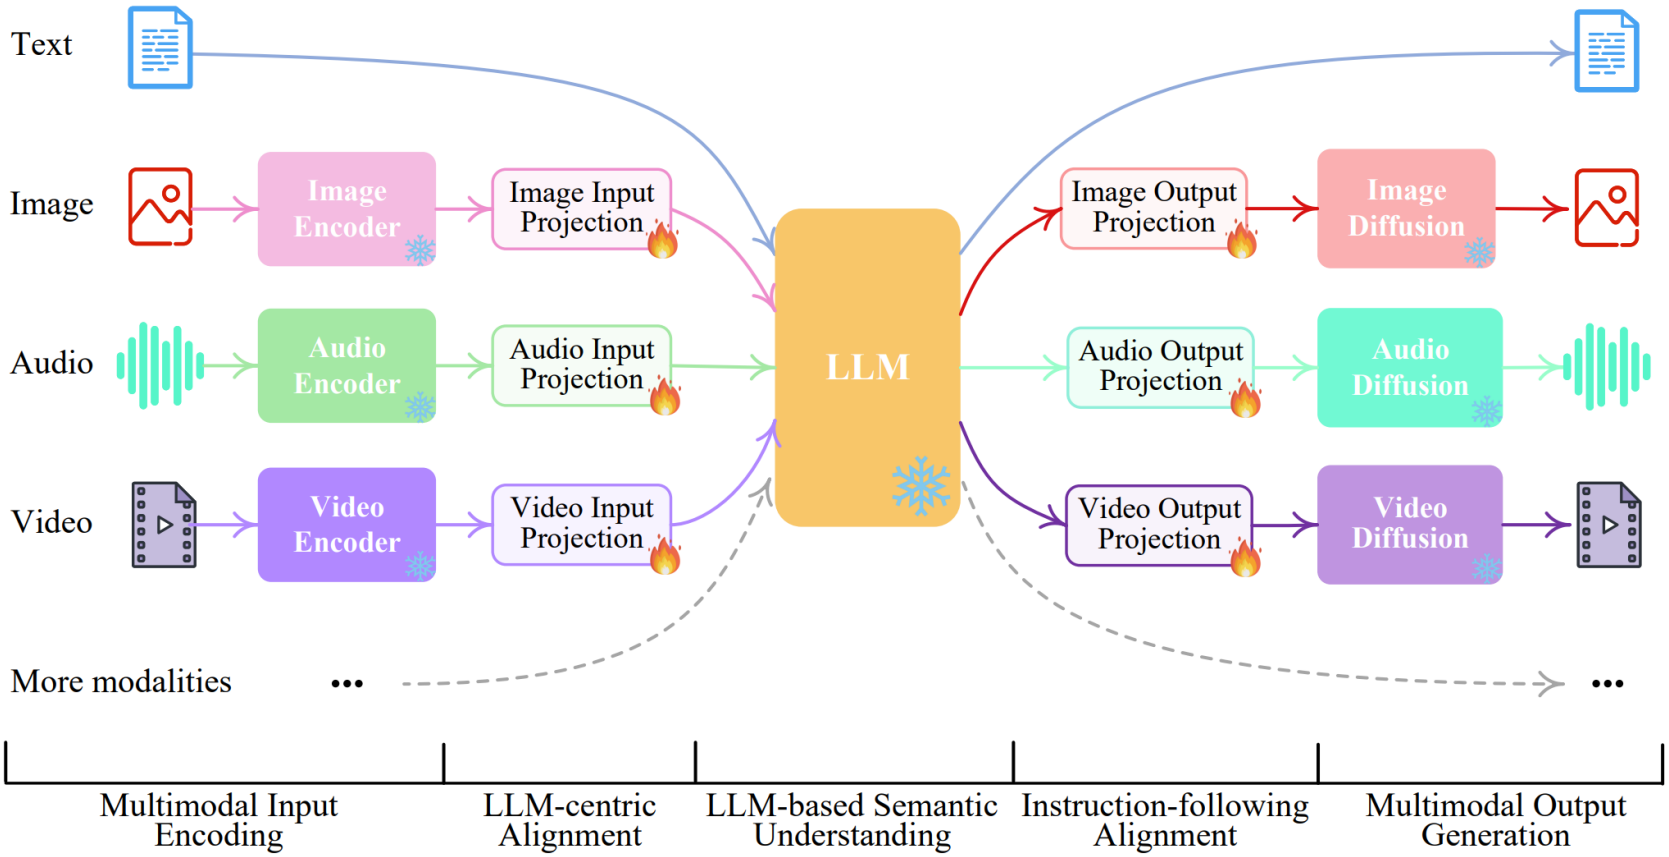

Wu, Shengqiong, et al. "Next-gpt: Any-to-any multimodal llm." Forty-first International Conference on Machine Learning. 2024.

LLM

- Limitations
  - Produces hallucinations
  - Lacks sufficient controllability
  - High training and inference costs
  - Difficulty maintaining context until maximum length
  - Struggles to update rapidly changing knowledge (temporal, regional, cultural)
  - Results difficult to interpret or verify without sources
  - Risk of sensitive information exposure (confidential data, PII)

Agent

- Definition
  - LLM Agents are autonomous systems that use large language models as their reasoning engine to understand tasks, make decisions, and take actions through external tools to achieve user-defined goals

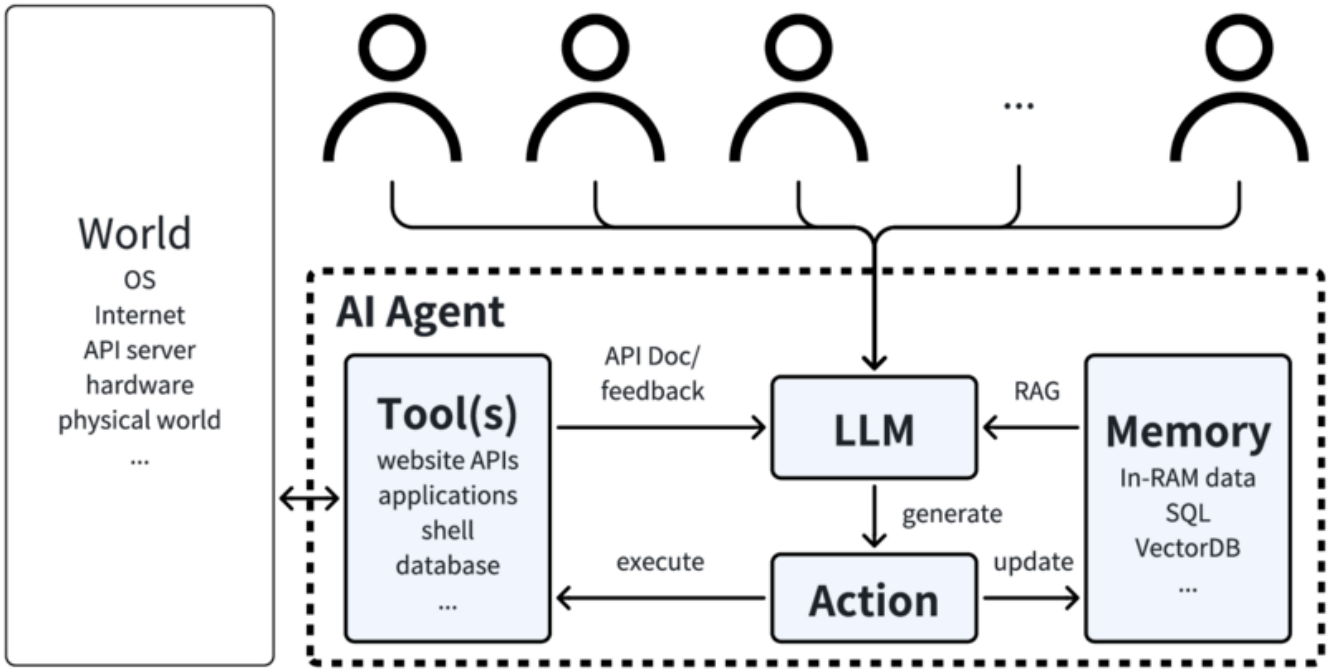

He, Yifeng, et al. "Security of ai agents." arXiv preprint arXiv:2406.08689 (2024).

Agent

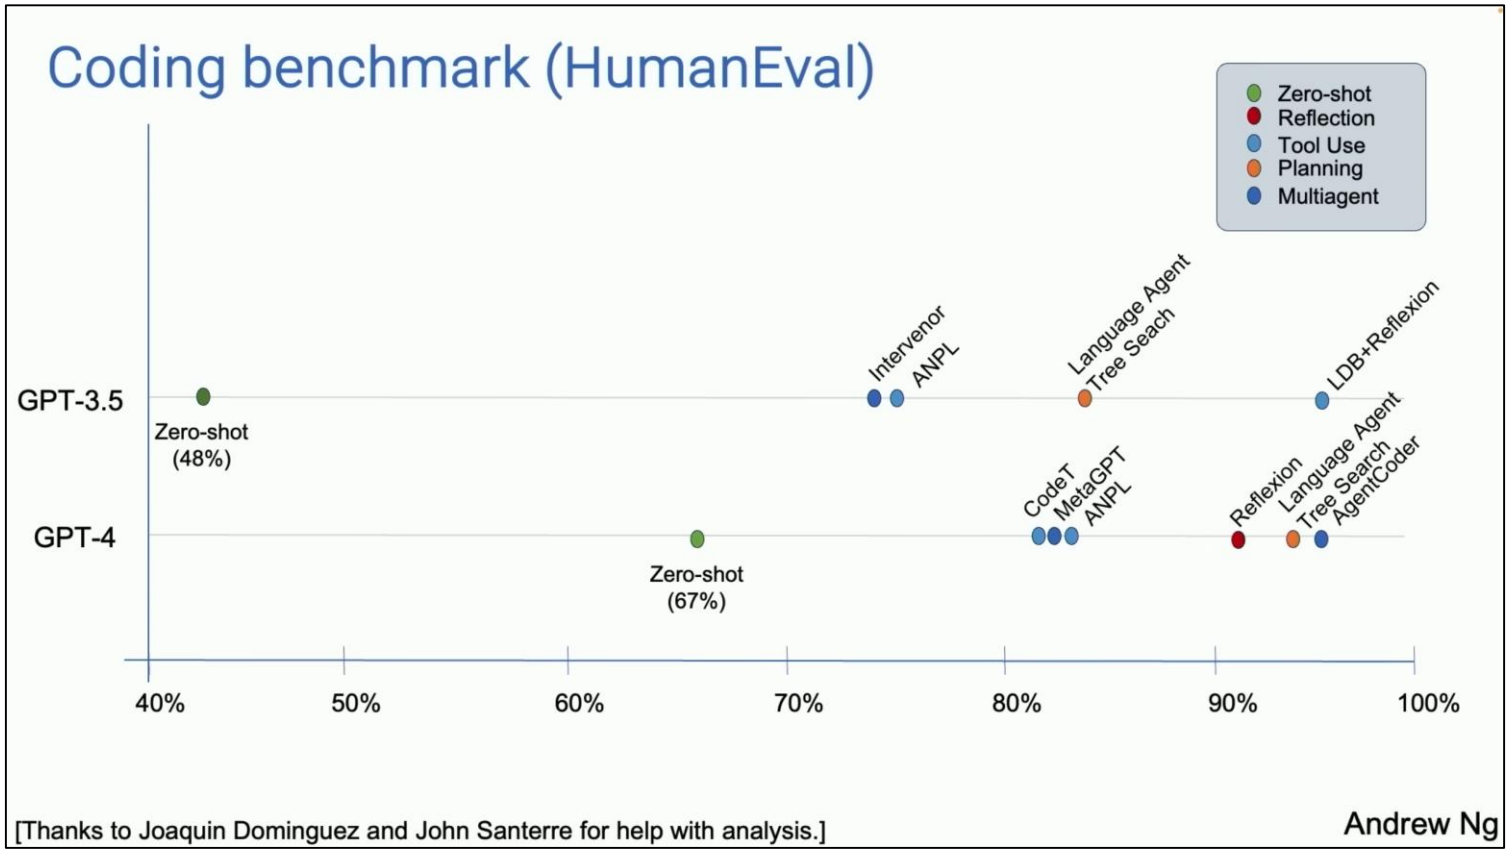

What's next for AI agentic workflows, https://youtu.be/sal78ACtGTc?si=zA7W9TnJekIBZsW2

RAG

- Definition
  - RAG combines external information retrieval with language model generation to produce more accurate and knowledge-grounded responses

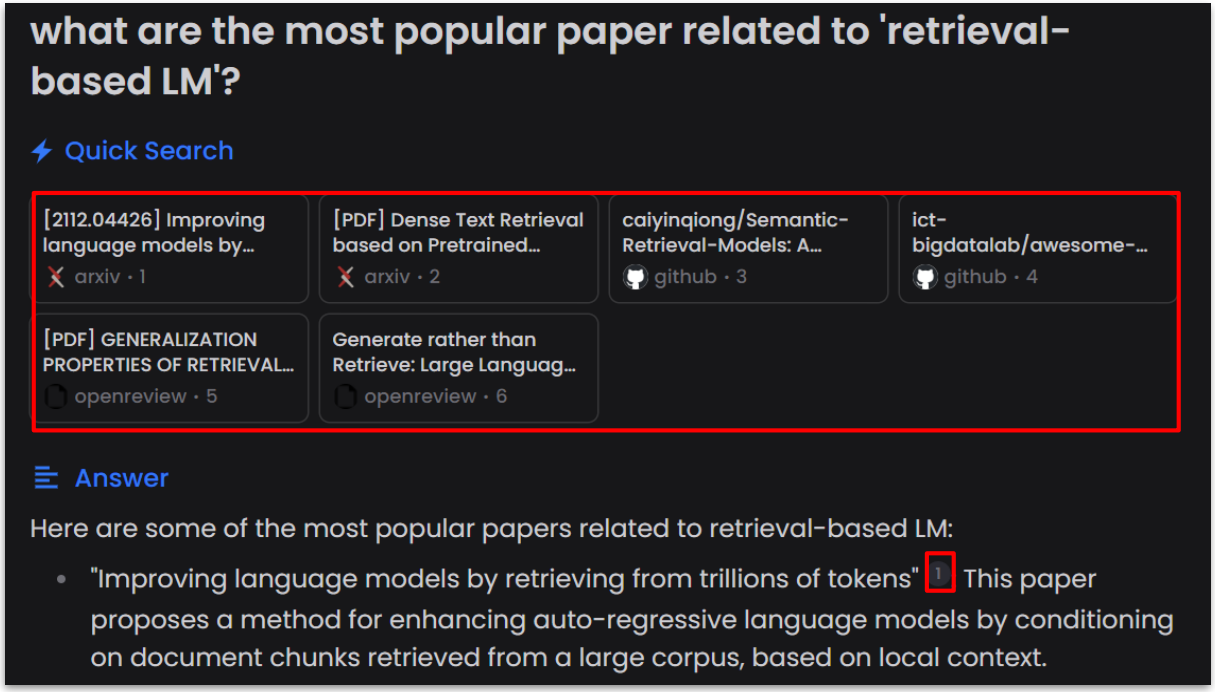

Knowledge Types

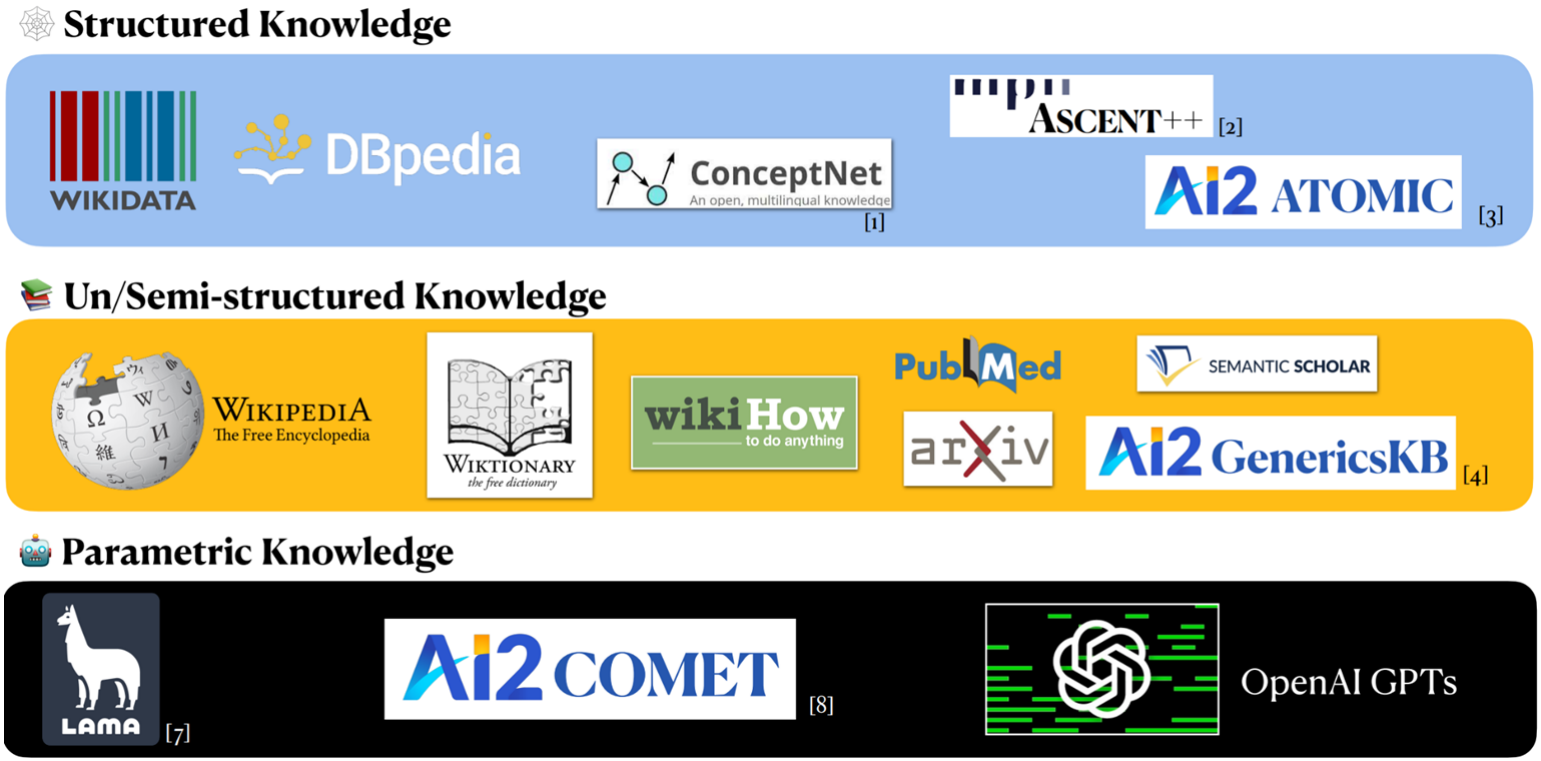

Knowledge Types

- Structured Knowledge
  - Vulnerability databases (CVE records with standardized fields)
  - MITRE ATT&CK framework tactics and techniques
  - Security event logs with defined fields and relationships
- Un/Semi-structured Knowledge
  - Threat intelligence reports and security blogs
  - Malware analysis narratives and researcher notes
  - Social media posts discussing emerging threats
- Parametric Knowledge
  - Microsoft Security Copilot
  - Google Sec-Gemini

RAG with Different Knowledge Types

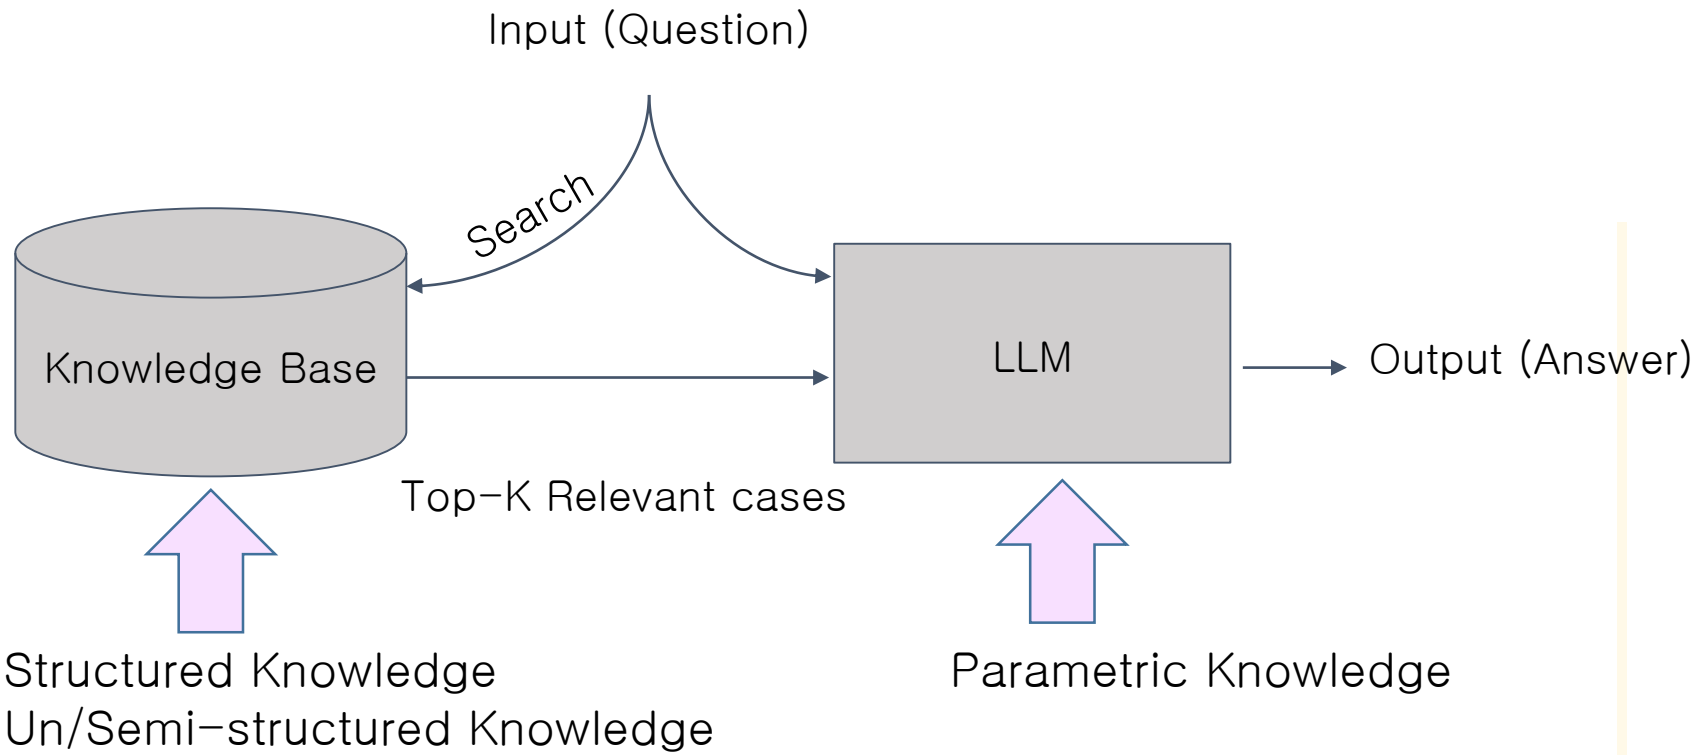

Other Terms

They are fascinating topics worth exploring in more depth, but for this class, just remember that modern LLMs can understand human instructions

https://www.thecloudgirl.dev/blog/top-15-llm-terms-you-need-to-know-in-2024

Generative AI in Cybersecurity

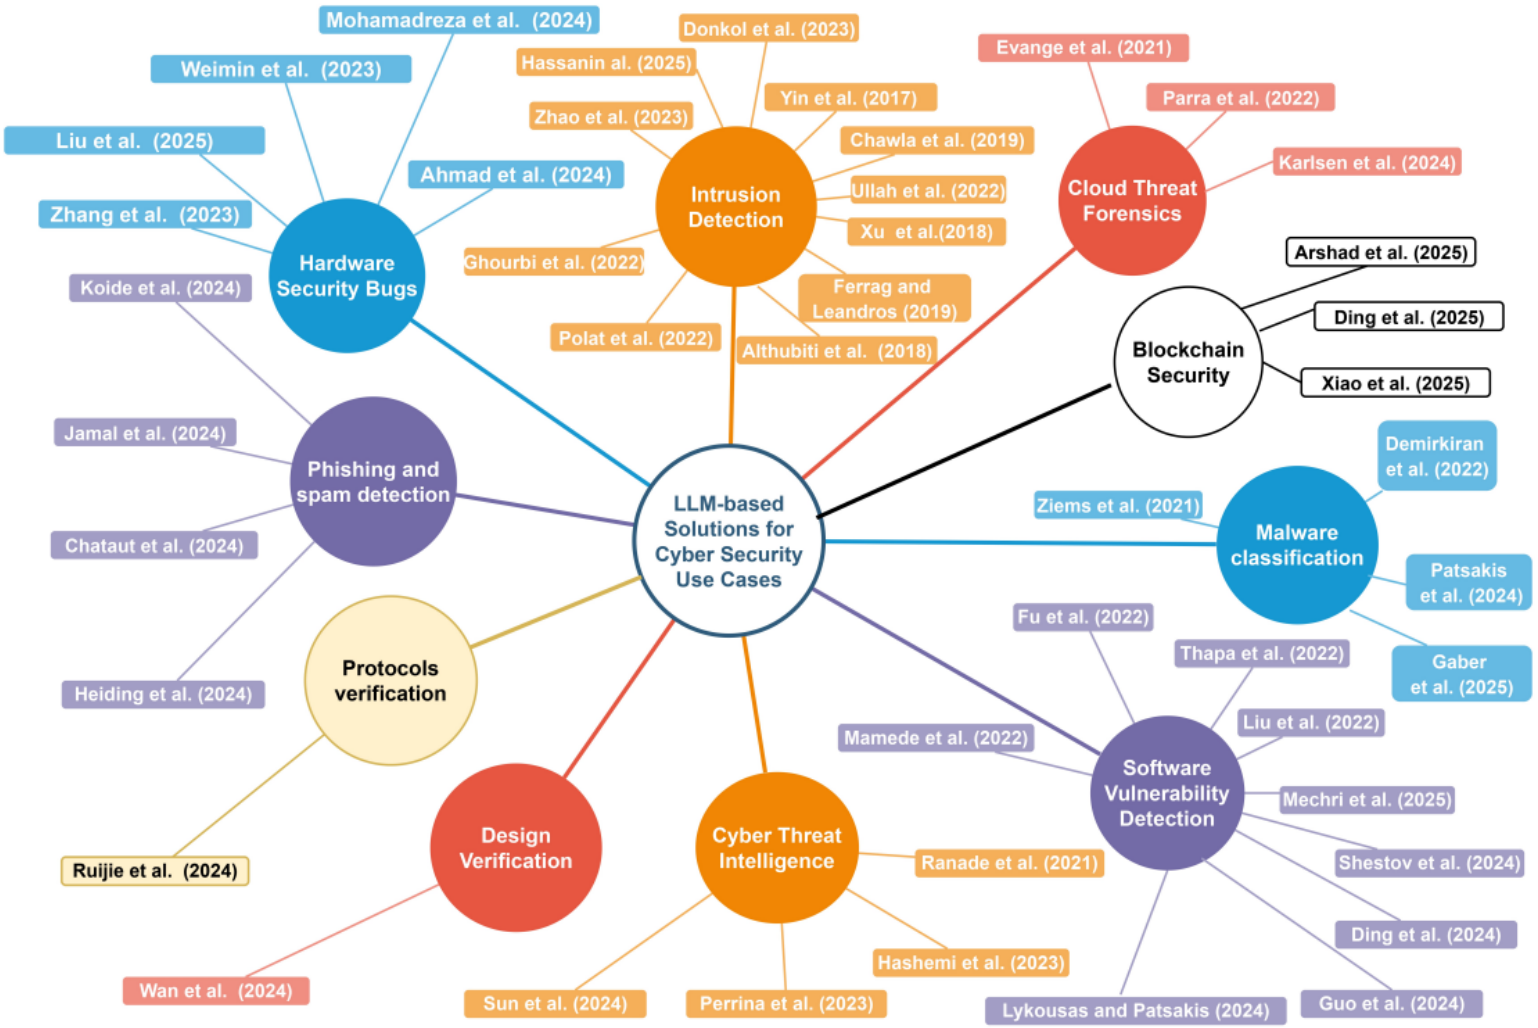

Ferrag, Mohamed Amine, et al. "Generative AI in Cybersecurity: A Comprehensive Review of LLM Applications and Vulnerabilities." Internet of Things and Cyber-Physical Systems (2025).

Use Case 1: MITRE ATT&CK Mapping

- Objective
  - AI tool can automate threat report analysis by mapping human language descriptions to MITRE ATT&CK frameworks without STIX/TAXII expertise
  - Five-step workflow: summarize materials, parse behaviors, map to techniques, verify results, and generate reports

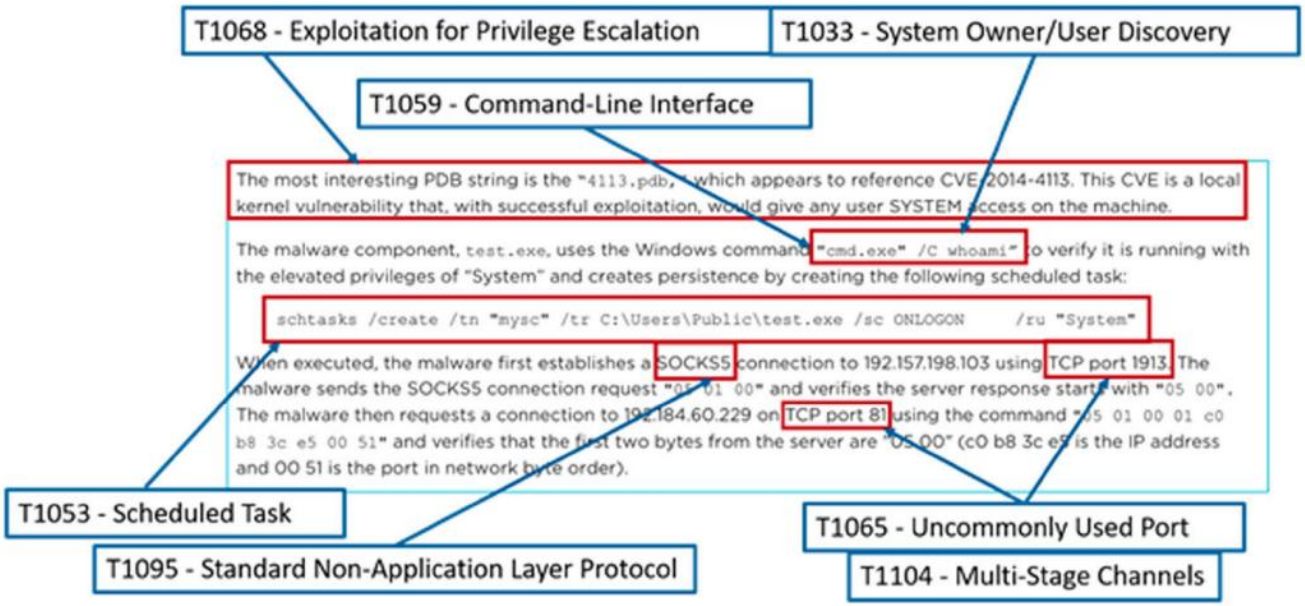

https://www.mitre.org/sites/default/files/2021-11/getting-started-with-attack-october-2019.pdf

Use Case 1: MITRE ATT&CK Mapping

- MITRE CWE and ATT&CK
  - CWE is a community-developed list of common software and hardware weaknesses, faults, and vulnerabilities
  - ATT&CK is a knowledge base maintained by MITRE that documents the tactics, techniques, and procedures (TTPs) used by adversaries during cyberattacks

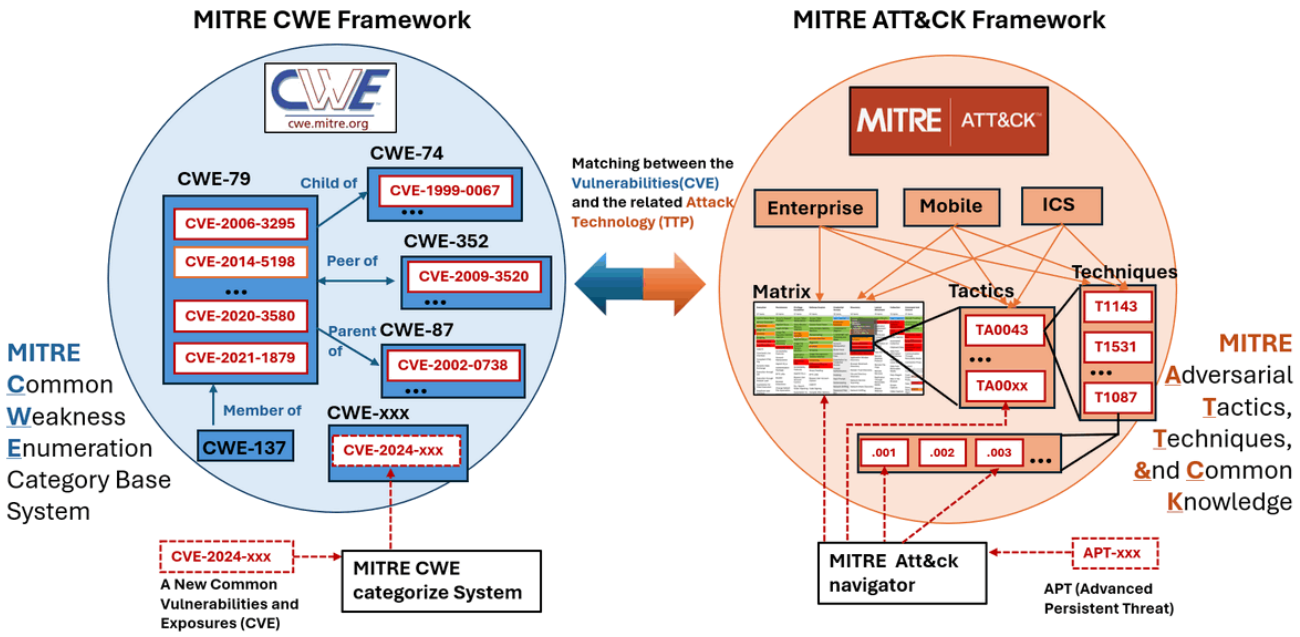

https://github.com/LiuYuancheng/Threats_2_MITRE_AI_Mapper

Use Case 1: MITRE ATT&CK Mapping

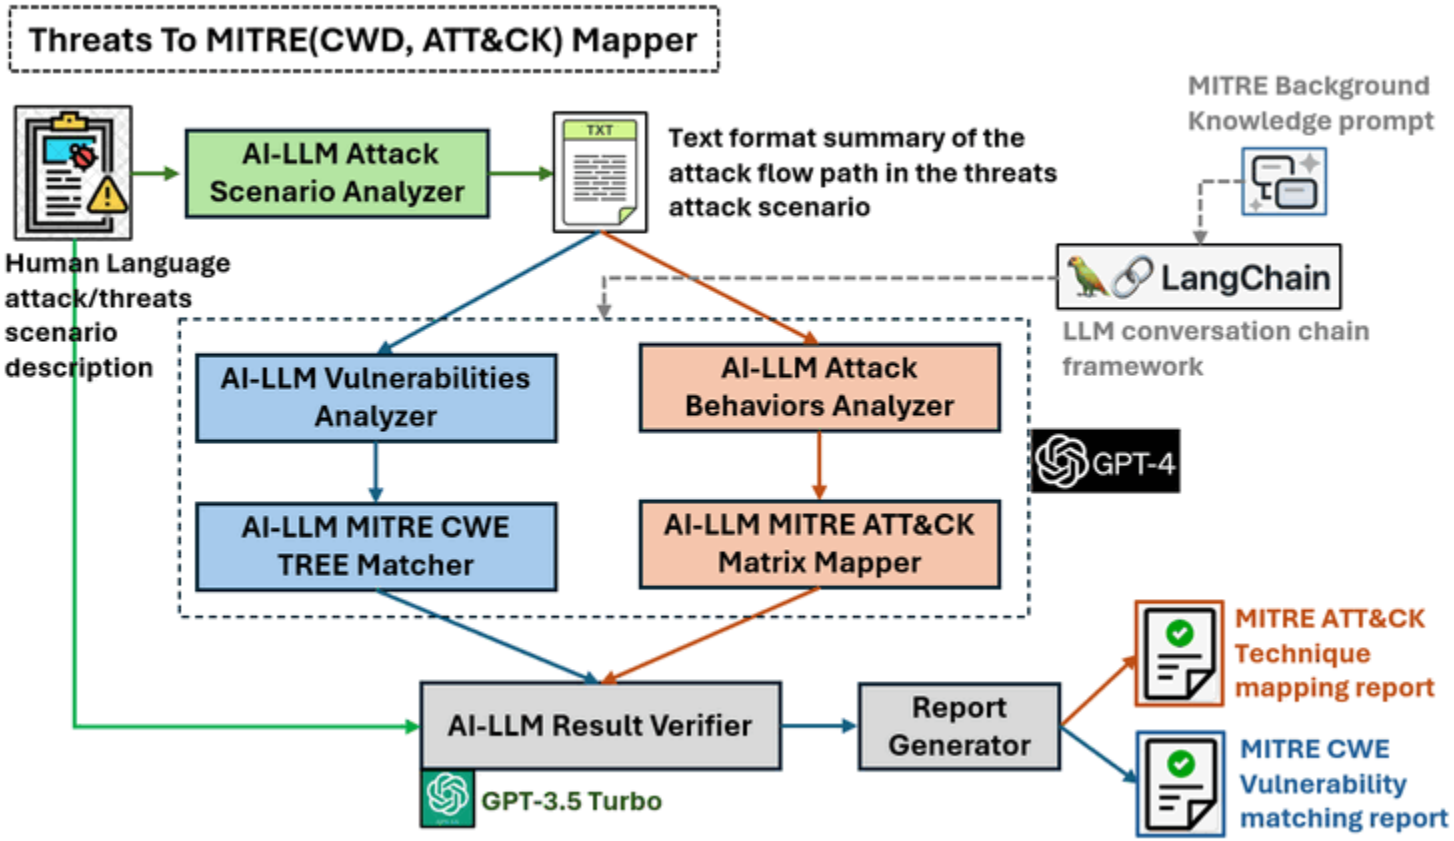

https://github.com/LiuYuancheng/Threats_2_MITRE_AI_Mapper

Use Case 1: MITRE ATT&CK Mapping

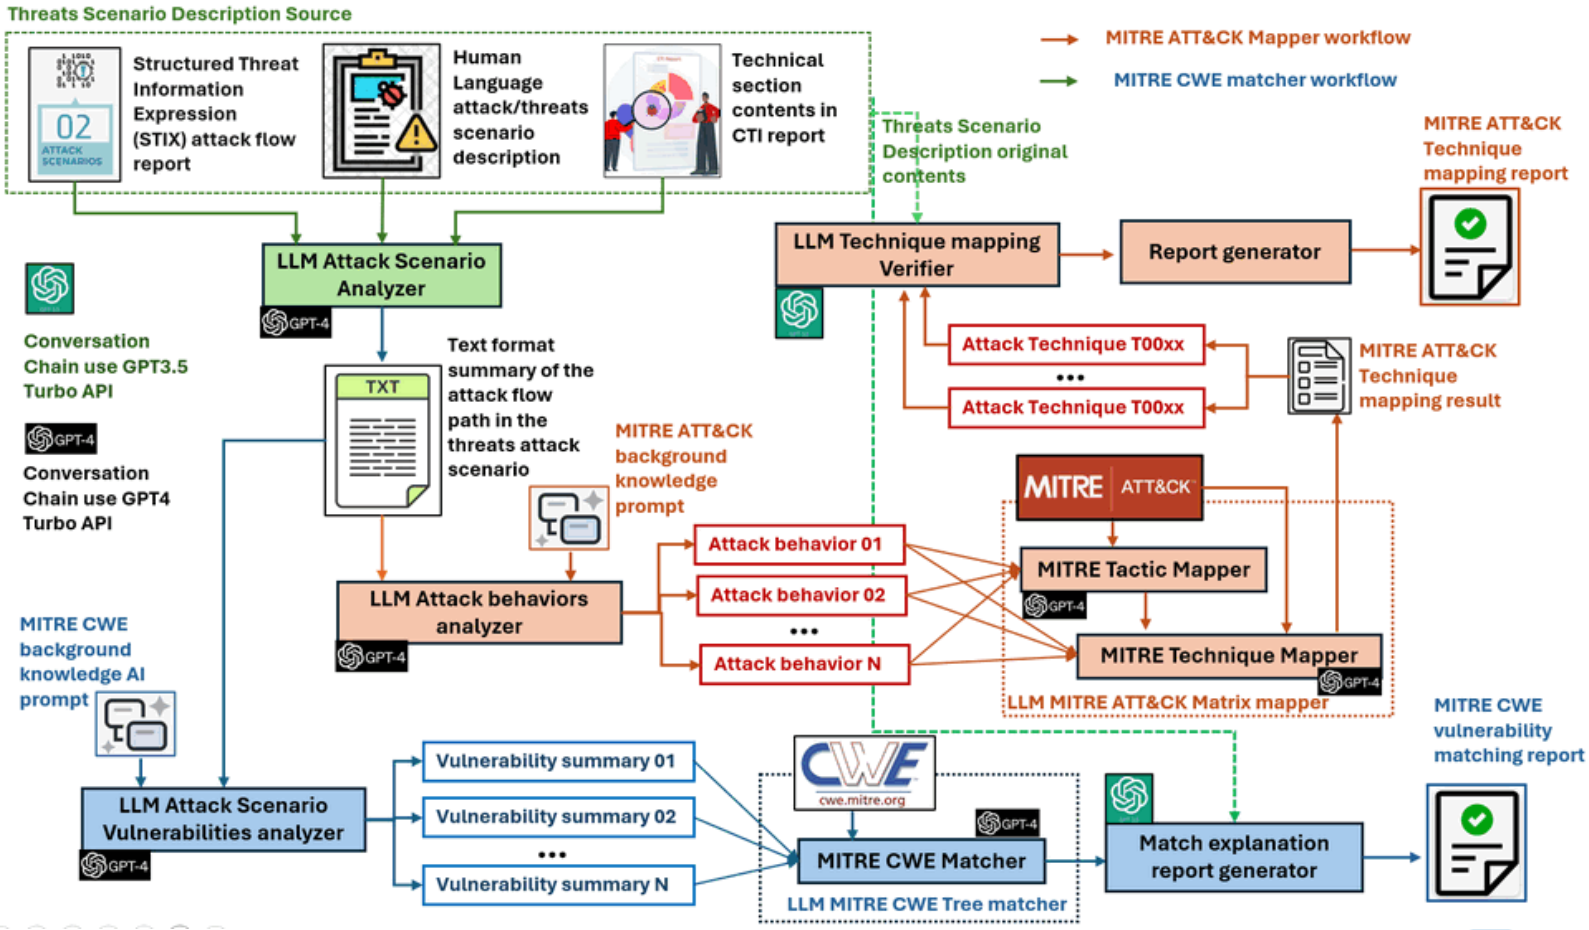

https://github.com/LiuYuancheng/Threats_2_MITRE_AI_Mapper

Use Case 2: Mapping SIEM Rules to TTPs

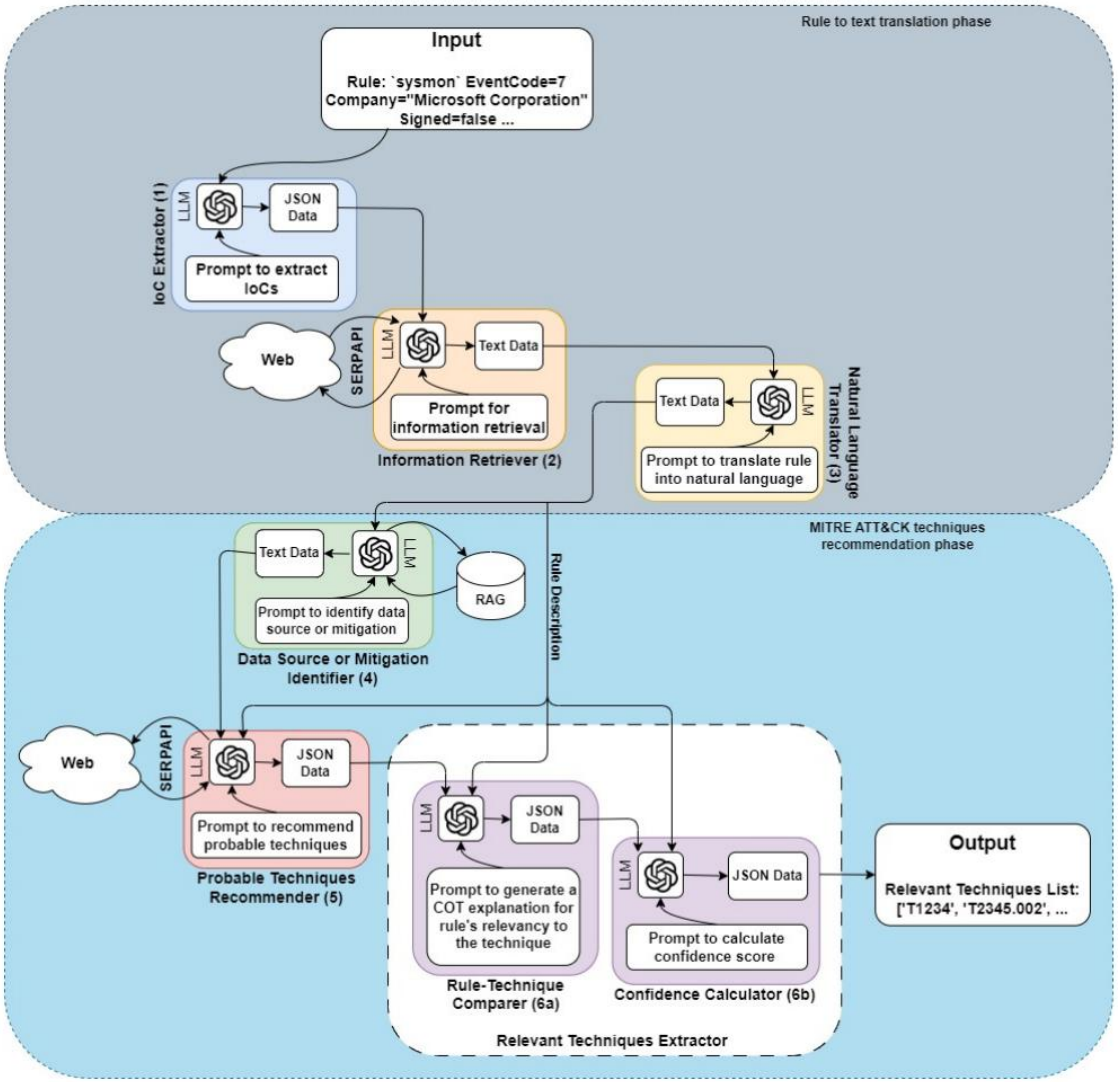

Wudali, Prasanna N., et al. "Rule-ATT&CK Mapper (RAM): Mapping SIEM Rules to TTPs Using LLMs." arXiv preprint arXiv:2502.02337 (2025).

(Hands-on 1) Data Collection and Processing

- Objectives
  - Learn how to collect and process threat intelligence data from various sources such as webpages, RSS feeds, and DNS records
  - Practice extracting indicators of compromise (IOCs) like IPs, domains, hashes, and filenames using web scraping and regular expressions
  - Explore feed parsing and DNS record resolution to enrich the context of security events
  - Gain experience with IP attribution by mapping IP addresses to autonomous systems using WHOIS data
- Python Packages
  - requests, beatifulsoup4, re, feedparser, dnspython, ipwhois

(Hands-on 2) Prompt Engineering for LLM

- Objectives
  - Learn how to integrate OpenAI's GPT models for automated classification of cyber attack techniques based on MITRE ATT&CK
  - Understand how to validate structured outputs using Pydantic models
  - Practice extracting and translating text from images using multi-modal prompts
- Python Packages
  - pydantic, openai

(Hands-on 3) Embedding for Similarity Search

- Objectives
  - Learn to apply embedding models to transform textual threat group descriptions into numerical vectors
  - Explore dimensionality reduction and visualization techniques using PCA
  - Analyze and compare threat actor similarities using cosine similarity and query matching
- Python Packages
  - requests, sentence_transformers, stix2, pandas, matplotlib, sklearn

(Hands-on 4) Querying Data with LLM

- Objectives
  - Build a natural language interface for pandas DataFrames using LangChain agents
  - Generate SQL queries from natural language questions using SQLCoder, and explore how LLMs can automate structured data access
- Python Packages
  - kagglehub, langchain, transformers, sqlparse

(Hands-on 5) Multi Agents based Threat Research Team

- Objectives
  - Build an automated cybersecurity investigation system using multiple specialized AI agents working together to analyze security alerts
  - Demonstrate how different security roles (SOC Analyst, Threat Intelligence, Reverse Engineering, and Phishing Analysis) can collaborate through a coordinated workflow
- Python Packages
  - autogen

Threat Research Team

SOC Analyst

Threat Intelligence Analyst

Phishing Analyst

Reverse Engineering Analyst

Task

Investigate a suspicious login alert

Investigate a malware detection alert

Investigate a data exfiltration alert

Investigate a phishing campaign

Thank you ☺